




Priya Nambiar
Senior Data Scientist

Priya Nambiar

Senior Data Scientist · ML · Statistical Modelling · MLOps

✉ priya.nambiar@datasci.io
☎ +91 98200 33481
📍 Bangalore, India
🐙 github.com/priya-nambiar-ds
🌐 linkedin.com/in/priyanambiar
🏆 kaggle.com/priyanambiar

TECHNICAL SKILLS



TOOLING

PyTorch / TensorFlow
XGBoost / LightGBM
Airflow / Prefect
dbt / Great Expectations
SageMaker / Vertex AI
Docker / Kubernetes
Plotly / Dash / Streamlit
SHAP / LIME (XAI)

CERTIFICATIONS

- Google Professional ML Engineer — 2024
- AWS MLS-C01 — 2022
- Databricks Associate — 2023
- Deep Learning Specialisation — Coursera 2020

LANGUAGES

English — Fluent (C2)
Hindi — Native
Tamil — Native
Kannada — Conversational

PUBLICATIONS

"Causal Inference in Churn Modelling" — KDD 2023
"Fairness-aware Recommendation" — RecSys 2022
2 patents pending (Flipkart)

Senior Data Scientist with 7 years transforming messy real-world data into revenue-driving products across e-commerce, fintech, and health-tech. Deep expertise in predictive modelling, causal inference, and operationalising models at scale via MLOps platforms. Comfortable leading cross-functional squads and presenting findings to C-suite stakeholders.

PROFESSIONAL EXPERIENCE

Senior Data Scientist

Flipkart · Supply Chain & Demand Forecasting

- Developed LightGBM demand-forecast model across 28M SKUs; reduced overstock costs by ₹180Cr annually.
- Built causal uplift model for coupon targeting — lifted incremental GMV by 12% vs prior propensity-score approach.
- Led MLOps migration to Vertex AI + MLflow; cut model deployment time from 3 weeks to 2 days.
- Designed A/B testing framework (CUPED) used across 40+ weekly experiments; trained 8 analysts.
- Mentored 3 junior data scientists; ran fortnightly internal reading group on ML fundamentals.

2022 - Present
Bangalore, IN

Data Scientist

Persistent Systems · AI Practice (BFSI Clients)

- Built credit risk scorecard (logistic regression + WoE binning) for ICICI Bank; Gini 0.68, deployed to 4M customers.
- Created NLP-based ticket triage model (BERT fine-tuned) reducing MTTR for L1 support by 35%.
- Delivered churn prediction pipeline for a telecom client using Spark + Airflow; 82% precision @top-decile.

2019 - 2022
Pune, IN

Data Analyst → Junior Data Scientist

Zoho Corporation · Analytics Platform

- Developed cohort analysis and funnel dashboards in Plotly/Dash used daily by 200+ product managers.
- First ML model shipped: anomaly detection on usage logs (Isolation Forest), catching billing fraud saving \$80K/yr.

2017 - 2019
Chennai, IN

NOTABLE PROJECTS

Price Elasticity Engine

Hierarchical Bayesian model estimating category-level price elasticity across 500 product categories with uncertainty bounds.

PyMC Stan Bayesian

Explainable Risk Dashboard

Streamlit app exposing SHAP waterfall charts to credit analysts — reduced model review time from 2 days to 4 hours.

SHAP Streamlit XGBoost

AutoML Experiment (2020)

H2O AutoML pipeline — abandoned due to poor explainability for regulated BFSI use case. Reverted to custom models.

H2O Deprecated

Open Source: faireval

Python library for fairness metrics (demographic parity, equalized odds). 680 PyPI downloads/month.

Python PyPI Ethics

EDUCATION

M.Tech. Data Science & Engineering

IIT Madras — Chennai, India

2015 - 2017 · CGPA 9.2 / 10 · Thesis: Bayesian Online Learning for Non-stationary Streams

B.Tech. Computer Science

NIT Trichy — Tiruchirappalli, India

2011 - 2015 · CGPA 8.8 / 10